

Programming

Problem 1:

VAE Training Parameter Details

1. Encoder structure

Four Linear layers of the following input and output size

- `nn.Linear(200, 128)`
- `nn.Linear(128, 64)`
- `nn.Linear(64, 32)`
- `nn.Linear(32, 16)`

Activation layer: ReLU

Batchnorm1d between linear layers

number of trainable parameters: 24768

2. Decoder structure

Four Linear layers of the following input and output size

- `nn.Linear(16, 32)`
- `nn.Linear(32, 64)`
- `nn.Linear(64, 128)`
- `nn.Linear(128, 200)`

Activation layer: ReLU, at the end of decoder tanh is used

Batchnorm1d between linear layers

number of trainable parameters: 24324

3. Learning rate: 0.001

4. Batch size: 512

5. Number of epochs: 2000


6. Alpha for kl divergence loss weightage = 0.005

7. Loss function = `MSELoss()` + `alpha*KL_divergence`, {reduction used is sum}

8. Optimizer = `Adam()`


9. Learning rate scheduler = `StepLR(step_size=100, gamma=0.4)`

Results




Real airfoils

Reconstructed airfoils



Synthesized airfoils



Training loss vs epoch

GAN

Problem 1:

GAN Training Parameter Details

10. Generator structure

Four Linear layers of the following input and output size

- Linear(in_features=16, out_features=32, bias=True)
- Linear(in_features=32, out_features=64, bias=True)
- Linear(in_features=64, out_features=128, bias=True)
- Linear(in_features=128, out_features=200, bias=True)

Activation layer: ReLU

number of trainable parameters: 36776

11. Discriminator structure

Four Linear layers of the following input and output size

- Linear(in_features=200, out_features=128, bias=True)
- Linear(in_features=128, out_features=64, bias=True)
- Linear(in_features=64, out_features=32, bias=True)
- Linear(in_features=32, out_features=1, bias=True)

Activation layer: ReLU,

number of trainable parameters: 36097

12. Gen Learning rate: 0.00015

13. Disc learning rate: 0.0002


14. Batch size: 512

15. Number of epochs: 5000


16. Loss function = BCEwithlogitsloss()

17. Optimizer = Adam() {for both gen and dis}

Gen loss
tag: Gen loss




Disc loss
tag: Disc loss



Loss vs number of iterations

Generated data from noise



Observation: The samples from gan are very noisy. The reason might be unlike VAE we do not have a reconstruction component in the loss, which compares each point generated individually in my code `MSELoss`. Also, VAE directly optimizes the ELBO function of MLE, while GAN is just classifying generated and real data. For gan, all we have is a classification loss which suggests if the data is fake or real. That being said the data from gans appear more diverse.